Alex Wang

AI & Cloud Solutions Engineer · NVIDIA-Stack & Azure AI Infrastructure

Los Angeles · San Francisco Bay Area   linkedin.com/in/alex-wang-1005oct   UCLA · Computer Science

PROFESSIONAL SUMMARY

AI & cloud-solutions engineer building the compute layer for the agentic-AI era at **Microsoft**. Practitioner experience spans **NVIDIA H100 / H200 / Blackwell GPU clusters**, **Azure AI** + **OpenAI-on-Azure** deployments, distributed training, and inference-serving topology. End-to-end fluency from silicon-side integration (NVIDIA partner ecosystem) to hyperscaler-side enterprise architecture. Active interest in **agentic AI**, **inference-time compute scaling**, and **hardware-aware ML frameworks**.

EXPERIENCE

AI & Cloud Solutions Engineer · Microsoft — *2025 — Present · United States*

Microsoft AI & Azure · Scaling AI Team

- Solutions engineering for enterprise AI workloads on **Azure AI**, **Azure OpenAI Service**, and hybrid on-prem topologies serving Fortune-500 customers.
- Architect AI training and inference clusters across **NVIDIA H100 / H200 / Blackwell** GPU pools with **InfiniBand** / **NVLink** fabric, NCCL-optimized collectives, CUDA-aware MPI.
- Bridge between Microsoft research, partner silicon (NVIDIA, AMD), and enterprise buyers — frontier model deployment into production architecture.
- Active work on **agentic AI** orchestration, tool use, and the systems engineering of long-horizon agent execution. Azure AI Fundamentals + NVIDIA DGX Cloud certified.

Solutions Engineer · AMAX Engineering — NVIDIA Partner Ecosystem — *Prior · United States*

GPU Systems & AI Infrastructure Integrator

- Designed and integrated **multi-GPU server topologies** for AI training and inference: NVLink / NVSwitch fabric, InfiniBand vs. RoCE trade-offs, PCIe Gen5 lane planning, liquid-cooled rack-level architecture.
- Customer-facing technical lead on AI training cluster deployments — capacity sizing, distributed-training framework selection (DeepSpeed, FSDP, Megatron-LM), inference-serving topology (vLLM, TensorRT-LLM, Triton).
- Operated inside the **NVIDIA partner ecosystem** — direct exposure to DGX / HGX / OEM reference platforms; certified NVIDIA solutions training.

EDUCATION

University of California, Los Angeles (UCLA) · B.S. Computer Science — *Los Angeles, CA*

Systems, algorithms, distributed computing, computer architecture. UCLA CS — feeder program to FAANG and frontier AI research labs.

CORE COMPETENCIES & STACK

AI Infrastructure: NVIDIA H100 · H200 · B200/Blackwell · GH200 Grace Hopper · DGX · HGX · NVLink · NVSwitch · InfiniBand (HDR/NDR) · RoCE v2 · CUDA · cuDNN · NCCL · RDMA · GPUDirect.

Frameworks & Serving: PyTorch · JAX · DeepSpeed · FSDP · Megatron-LM · vLLM · TensorRT-LLM · Triton Inference Server · Ray · Hugging Face Transformers.

Cloud & Platforms: Microsoft Azure · Azure AI · Azure OpenAI Service · Azure ML · AKS · AWS · GCP · Kubernetes · Docker · Terraform.

Agentic & LLM Apps: LangChain · LangGraph · LlamaIndex · ReAct · MCP (Model Context Protocol) · function calling · tool use · RAG · multi-agent orchestration.

Languages: Python · C++ · CUDA C · Go · TypeScript · Bash.   **Certs:** Microsoft Azure AI Fundamentals · NVIDIA DGX Cloud Solutions.

Focus areas: inference-time compute scaling · agentic AI systems · hardware-aware ML · frontier-lab research (Anthropic / OpenAI / Google DeepMind) · enterprise AI GTM.

Keywords: AI infrastructure engineer, ML systems, ML platform, AI engineer, NVIDIA GPU, H100, H200, B200/Blackwell, GH200, DGX, HGX, CUDA, NCCL, NVLink, InfiniBand, RDMA, GPUDirect, distributed training, model parallelism, tensor parallelism, FSDP, DeepSpeed, Megatron, vLLM, TensorRT-LLM, Triton, inference optimization, KV cache, FP8 quantization, Microsoft, Azure AI, Azure OpenAI, OpenAI, Anthropic, Claude, Google DeepMind, Gemini, agentic AI, multi-agent systems, tool use, function calling, MCP, RAG, LangGraph, LangChain, ReAct, reasoning models, inference-time compute, foundation models, LLM, enterprise AI, AI solutions architect, MLOps, GPU cluster, hyperscaler, UCLA, frontier AI lab, GPU systems engineer.