

Alex Wang

AI & Cloud Solutions Engineer · NVIDIA-Stack & Azure AI Infrastructure

Los Angeles · San Francisco Bay Area [linkedin.com/in/alex-wang-1005oct](https://www.linkedin.com/in/alex-wang-1005oct) UCLA · Computer Science

PROFESSIONAL SUMMARY

AI & cloud-solutions engineer building the compute layer for the agentic-AI era at **Microsoft**. Practitioner experience across the modern AI stack: **NVIDIA H100 / H200 / Blackwell GPU clusters**, **Azure AI** and **OpenAI-on-Azure** deployments, distributed training and inference-serving topology. Rare end-to-end exposure spanning silicon-side integration (NVIDIA partner ecosystem) and hyperscaler-side enterprise architecture (Microsoft Azure AI). Active interest in **agentic AI systems**, **inference-time compute scaling**, and **hardware-aware ML frameworks**.

EXPERIENCE

AI & Cloud Solutions Engineer · Microsoft

2025 — Present

United States · Microsoft AI & Azure · Scaling AI Team

- Solutions engineering for enterprise AI workloads on **Azure AI**, **OpenAI on Azure**, and hybrid on-prem topologies serving Fortune-500 customers.
- Architect AI training and inference clusters across **NVIDIA H100 / H200 / Blackwell GPU** pools with **InfiniBand / NVLink** fabric, NCCL-optimized collectives, and CUDA-aware MPI.
- Bridge between Microsoft research, partner silicon (NVIDIA, AMD), and enterprise buyers — translating frontier model deployment requirements into production architecture.
- Active work on **agentic AI** orchestration patterns, tool-use frameworks, and the systems-engineering of long-horizon agent execution.
- Azure AI Fundamentals certified · NVIDIA DGX Cloud certified.

Solutions Engineer · AMAX Engineering — NVIDIA Partner Ecosystem

Prior

United States · GPU Systems & AI Infrastructure Integrator

- Designed and integrated **multi-GPU server topologies** for AI training and inference customers: **NVLink / NVSwitch** fabric design, InfiniBand vs. RoCE trade-offs, PCIe Gen5 lane planning, liquid-cooled rack-level architecture.
- Customer-facing technical lead on AI training cluster deployments — capacity sizing for foundation-model training, distributed-training framework selection (DeepSpeed, FSDP, Megatron-LM), and inference-serving topology (vLLM, TensorRT-LLM, Triton).
- Operated inside the **NVIDIA partner ecosystem** — direct exposure to DGX, HGX, and OEM reference platforms; certified NVIDIA solutions engineering training.

EDUCATION

University of California, Los Angeles (UCLA) · B.S. Computer Science

Los Angeles, CA

Systems, algorithms, distributed computing, computer architecture. UCLA CS is a feeder program to FAANG and frontier AI research labs.

CORE COMPETENCIES & STACK

AI Infrastructure: NVIDIA H100, H200, B200/Blackwell, GH200 Grace Hopper, DGX, HGX, NVLink, NVSwitch, InfiniBand (HDR/NDR), RoCE v2, CUDA, cuDNN, NCCL, RDMA, GPUDirect.

Frameworks & Serving: PyTorch, JAX, DeepSpeed, FSDP, Megatron-LM, vLLM, TensorRT-LLM, Triton Inference Server, Ray, Hugging Face Transformers.

Cloud & Platforms: Microsoft Azure, Azure AI, Azure OpenAI Service, Azure ML, Azure Kubernetes Service (AKS), AWS, GCP, Kubernetes, Docker, Terraform.

Agentic & LLM Apps: LangChain, LangGraph, LlamaIndex, ReAct, MCP (Model Context Protocol), function calling, tool use, RAG architectures, agent orchestration, multi-agent systems.

Languages: Python, C++, CUDA C, Go, TypeScript, Bash.

Certifications: Microsoft Azure AI Fundamentals · NVIDIA DGX Cloud Solutions.

AREAS OF FOCUS

- **Inference-time compute scaling** — the shift from training-FLOPs to inference-FLOPs and what it means for GPU-cluster economics.
- **Agentic AI systems** — long-horizon planning, tool use, multi-agent orchestration, autonomous task execution.
- **Hardware-aware ML** — memory hierarchy, kernel scheduling, NCCL behavior, mixed-precision training (FP8/FP16/BF16).
- **Frontier-lab research** — close reading of releases from Anthropic, OpenAI, Google DeepMind on reasoning, alignment, and agent architectures.
- **Enterprise AI GTM** — translating research velocity into CIO-procurement-cycle deployments at hyperscaler scale.

Keywords: AI infrastructure engineer, machine learning systems, ML platform, ML infra, AI engineer, AI/ML, NVIDIA GPU, H100, H200, Blackwell, B200, DGX, HGX, GH200, CUDA, cuDNN, NCCL, NVLink, NVSwitch, InfiniBand, RDMA, GPUDirect, distributed training, model parallelism, tensor parallelism, pipeline parallelism, FSDP, DeepSpeed, Megatron, vLLM, TensorRT-LLM, Triton, inference optimization, KV cache, quantization, FP8, mixed precision, Microsoft, Azure, Azure AI, Azure OpenAI, Azure ML, OpenAI, Anthropic, Claude, Google DeepMind, Gemini, agentic AI, agents, multi-agent systems, tool use, function calling, MCP, RAG, retrieval augmented generation, LangGraph, LangChain, LlamaIndex, ReAct, reasoning models, inference-time compute, test-time compute, foundation models, LLM, large language

model, enterprise AI, AI solutions architect, solutions engineer, AI/ML platform, ML ops, MLOps, AI infrastructure, GPU cluster, AI training cluster, AI compute, hyperscaler, AI workloads, UCLA, frontier AI lab, AI research engineer, ML systems engineer, GPU systems engineer.